

Linear cost of sequential ensemble inference
(vLLM multi-LoRA would amortize, ~2-3x predicted)

